

1. Neural Lattice Language Models (Buckman and Neubig 2018)
2. Quasi-Recurrent Neural Networks (Bradbury et al. 2016)
3. Learned in translation: Contextualized word vectors (McCann et al. 2017)
4. Deep contextualized word representations Blogpost (Peters et al. 2018)
5. Google Vizier: A Service for Black-Box Optimization (Golovin et al. 2017)

References

- Bradbury, James, Stephen Merity, Caiming Xiong, and Richard Socher. 2016. “Quasi-Recurrent Neural Networks.” *arXiv Preprint arXiv:1611.01576*.
- Buckman, Jacob, and Graham Neubig. 2018. “Neural Lattice Language Models.” *arXiv Preprint arXiv:1803.05071*.
- Golovin, Daniel, Benjamin Solnik, Subhodeep Moitra, Greg Kochanski, John Karro, and D Sculley. 2017. “Google Vizier: A Service for Black-Box Optimization.” In *Proceedings of the 23rd Acm Sigkdd International Conference on Knowledge Discovery and Data Mining*, 1487–95. ACM.
- McCann, Bryan, James Bradbury, Caiming Xiong, and Richard Socher. 2017. “Learned in Translation: Contextualized Word Vectors.” In *Advances in Neural Information Processing Systems*, 6297–6308.
- Peters, Matthew E, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. 2018. “Deep Contextualized Word Representations.” *arXiv Preprint arXiv:1802.05365*.